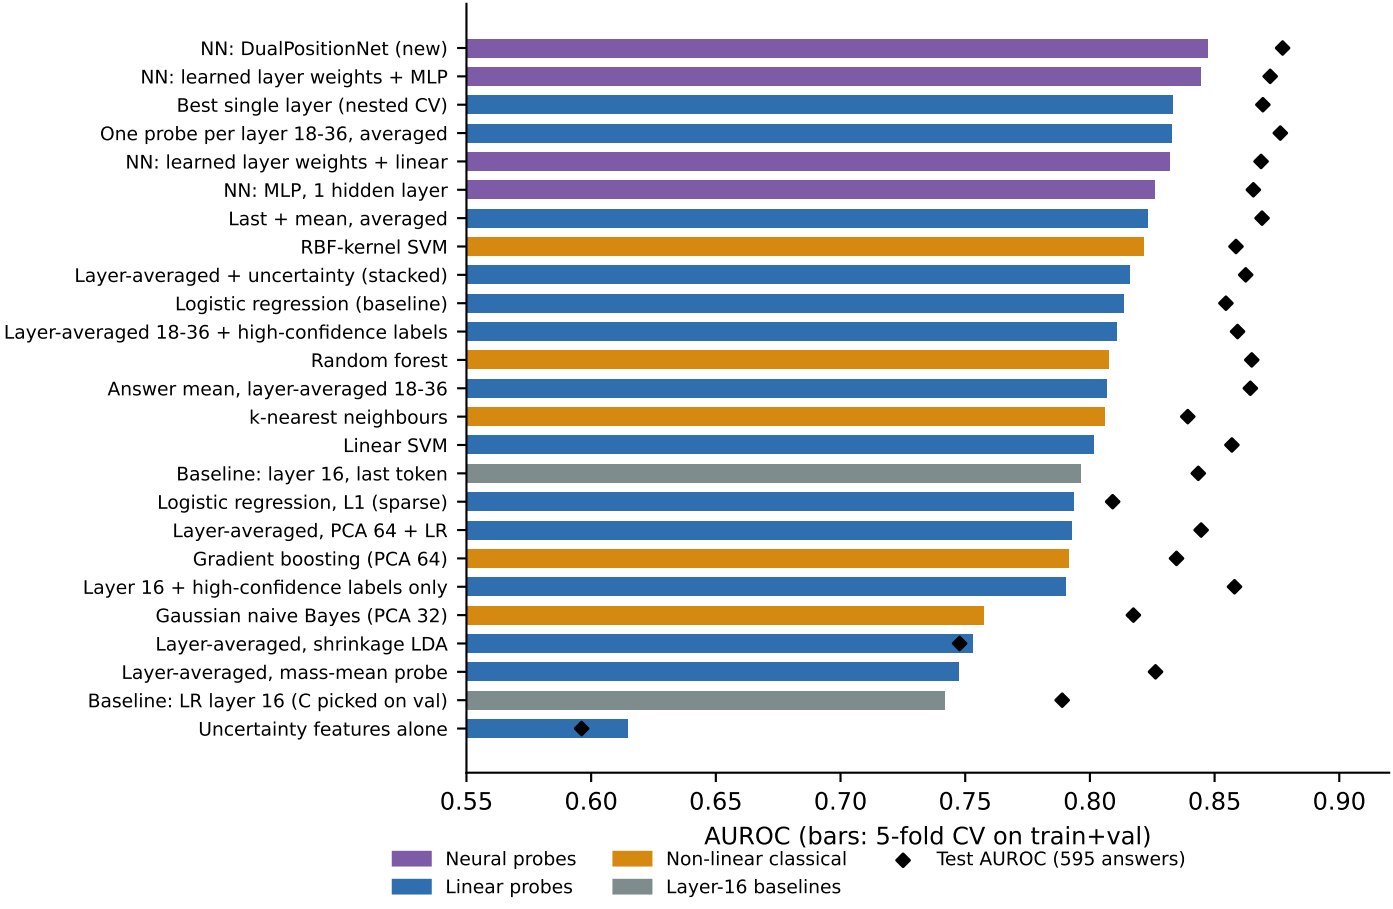
NN: DualPositionNet (new)
NN: learned layer weights + MLP
Best single layer (nested CV)
One probe per layer 18-36, averaged
NN: learned layer weights + linear
NN: MLP, 1 hidden layer
Last + mean, averaged
RBF-kernel SVM
Layer-averaged + uncertainty (stacked)
Logistic regression (baseline)
Layer-averaged 18-36 + high-confidence labels
Random forest
Answer mean, layer-averaged 18-36
k-nearest neighbours
Linear SVM
Baseline: layer 16, last token
Logistic regression, L1 (sparse)
Layer-averaged, PCA 64 + LR
Gradient boosting (PCA 64)
Layer 16 + high-confidence labels only
Gaussian naive Bayes (PCA 32)
Layer-averaged, shrinkage LDA
Layer-averaged, mass-mean probe
Baseline: LR layer 16 (C picked on val)
Uncertainty features alone
0.55
0.60
0.65
0.70
0.75
0.80
0.85
0.90
AUROC (bars: 5-fold CV on train+val)
Neural probes
Linear probes
Non-linear classical
Layer-16 baselines
Test AUROC (595 answers)